

1. Train a model to predict the fare of a taxi ride in Chicago, Illinois with the dataset, City of Chicago Taxi Trips dataset.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import kagglehub
import os
```

```
path = kagglehub.dataset_download(
    "adelanseau/taxi-trips-chicago-2024"
)

print("Path to dataset files:", path)
```

```
Downloading to /root/.cache/kagglehub/datasets/adelanseau/taxi-trips-chicago-2024/1.archive...
100%|██████████| 111M/111M [00:01<00:00, 89.5MB/s] Extracting files...

Path to dataset files: /root/.cache/kagglehub/datasets/adelanseau/taxi-trips-chicago-2024/versions/1
```

```
print(os.listdir(path))

['Taxi_Trips_-_2024_20240408.csv']
```

```
data = pd.read_csv(
    f"{path}/Taxi_Trips_-_2024_20240408.csv"
)

data.head()
```

	Trip ID	Taxi ID	Trip Start Timestamp	Trip End Timestamp	Trip Seconds
0	0287f53fedcde6b0af9aab9e32cbd2cadb337eaa	e54db25f18193a08f1f5754515e8c338480e04fb938ed3...	03/01/2024 12:00:00 AM	03/01/2024 12:00:00 AM	15.0
1	1f0034299b914fd7b45002568576a353c83acb2f	4a263f78f56815be4acd38658af8fc1824ce4e15f7ec81...	03/01/2024 12:00:00 AM	03/01/2024 12:15:00 AM	900.0
2	01a0b77b722a0b91f45cb4fc90b3e64f76fd0681	4f1e94982e6851725add382f7981d64006ae7c38f3664f...	03/01/2024 12:00:00 AM	03/01/2024 12:00:00 AM	711.0
3	14aff2071ac81c6450a8f8b0e1534497da900197	de8e8659ceb9eb0da842a46b60c0a5207098ac69fc23ab...	03/01/2024 12:00:00 AM	03/01/2024 12:30:00 AM	1770.0
4	16d3a633e2f380cee4a7fa4677db163b6e91255d	d79d3e19a1d5f6cb0cd4449d5579412262a2f1b182799d...	03/01/2024 12:00:00 AM	03/01/2024 12:15:00 AM	849.0

5 rows × 23 columns

```
df = pd.read_csv(
    f"{path}/Taxi_Trips_-_2024_20240408.csv"
)
```

```
print("Number of rows:", df.shape[0])

print("Number of columns:", df.shape[1])

print("\nColumn names:")

print(df.columns.tolist())

Number of rows: 865247
Number of columns: 23

Column names:
['Trip ID', 'Taxi ID', 'Trip Start Timestamp', 'Trip End Timestamp', 'Trip Seconds', 'Trip Miles', 'Pickup Census Tract', 'Dropoff Census Tract', 'Trip Status', 'Trip Type', 'Trip Category', 'Trip Subcategory', 'Trip Subcategory2', 'Trip Subcategory3', 'Trip Subcategory4', 'Trip Subcategory5', 'Trip Subcategory6', 'Trip Subcategory7', 'Trip Subcategory8', 'Trip Subcategory9', 'Trip Subcategory10', 'Trip Subcategory11', 'Trip Subcategory12', 'Trip Subcategory13', 'Trip Subcategory14', 'Trip Subcategory15', 'Trip Subcategory16', 'Trip Subcategory17', 'Trip Subcategory18', 'Trip Subcategory19', 'Trip Subcategory20', 'Trip Subcategory21', 'Trip Subcategory22', 'Trip Subcategory23']
```

```
df.describe()
```

	Trip Seconds	Trip Miles	Pickup Census Tract	Dropoff Census Tract	Pickup Community Area	Dropoff Community Area	Fare	Tips	
count	865078.000000	865241.000000	3.272510e+05	3.138540e+05	842052.000000	781394.000000	863202.000000	863202.000000	86
mean	1150.532018	6.563330	1.703152e+10	1.703141e+10	36.660485	25.853590	22.057441	2.761349	
std	1584.243513	8.279528	3.721479e+05	3.392339e+05	26.416845	20.501911	25.145823	4.132953	
min	0.000000	0.000000	1.703101e+10	1.703101e+10	1.000000	1.000000	0.000000	0.000000	
25%	459.000000	0.910000	1.703108e+10	1.703108e+10	8.000000	8.000000	8.000000	0.000000	
50%	900.000000	3.070000	1.703132e+10	1.703132e+10	32.000000	28.000000	15.750000	0.000000	
75%	1591.000000	11.710000	1.703198e+10	1.703184e+10	69.000000	32.000000	33.750000	4.000000	
max	86340.000000	3093.470000	1.703198e+10	1.703198e+10	77.000000	77.000000	9999.750000	200.000000	

```
data = data[['Trip Miles', 'Trip Total']]
```

```
data.shape
```

```
(865247, 2)
```

```
data.dropna(inplace=True)
```

```
data.isnull().sum()
```

```

      0
Trip Miles  0
Trip Total  0

dtype: int64
```

```
x = data.iloc[:, :1].values
y = data.iloc[:, 1:2].values
```

```
print(x)
print(y)
```

```
[[ 0.09]
 [ 3.   ]
 [ 5.84]
 ...
 [17.44]
 [ 3.07]
 [ 1.49]]
[[38.75]
 [12.   ]
 [26.7  ]
 ...
 [66.3  ]
 [26.   ]
 [12.25]]
```

```
x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2, random_state=42
)
```

```
model = LinearRegression()

model.fit(x_train, y_train)

y_pred = model.predict(x_test)

y_pred
```

```
array([[42.1919655 ],
       [28.15969042],
       [12.27758623],
       ...,
       [ 9.81995614],
       [48.63994124],
       [48.11141863]])
```

```
model.score(x_test, y_test)
```

```
0.1719224457128684
```

```
print(y_test)
print(y_pred)
```


```
[[31.25]
 [24.98]
 [ 8.75]
 ...
 [45.25]
 [43.5 ]
 [50.4 ]]
[[42.1919655 ]
 [28.15969042]
 [12.27758623]
 ...
 [ 9.81995614]
 [48.63994124]
 [48.11141863]]
```

```
plt.scatter(x_train, y_train, color='blue')

plt.plot(
    x_train,
    model.predict(x_train),
    color='red'
)

plt.title('Trip Miles vs Trip Total (Training Set)')
plt.xlabel('Miles')
plt.ylabel('Total in dollars')

plt.show()
```




```
plt.scatter(x_test, y_test, color='blue')

plt.plot(
    x_test,
    model.predict(x_test),
    color='red'
)

plt.title('Trip Miles vs Trip Total (Testing Set)')
plt.xlabel('Miles')
plt.ylabel('Total in dollars')

plt.show()
```



2. Conduct Linear regression on a model that predicts the body weight of a person. The dependent variable is thus the body weight, while the height, age and gender are chosen as independent variables. The following example data set is available:

Weight	Height	Age	Gender
79	1.80	35	Male
69	1.68	39	Male
73	1.82	25	Male
95	1.70	60	Male
82	1.87	27	Male
55	1.55	18	Female
69	1.50	89	Female
71	1.78	42	Female
64	1.67	16	Female
69	1.64	52	Female

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
```

Enter the given dataset

```
data = pd.DataFrame({
    'Weight': [79, 69, 73, 95, 82, 55, 69, 71, 64, 69],
    'Height': [1.80, 1.68, 1.82, 1.70, 1.87, 1.55, 1.50, 1.78, 1.67, 1.64],
    'Age': [35, 39, 25, 60, 27, 18, 89, 42, 16, 52],
    'Gender': ['Male', 'Male', 'Male', 'Male', 'Male',
               'Female', 'Female', 'Female', 'Female', 'Female']
})

data
```

	Weight	Height	Age	Gender	
0	79	1.80	35	Male	 
1	69	1.68	39	Male	
2	73	1.82	25	Male	
3	95	1.70	60	Male	
4	82	1.87	27	Male	
5	55	1.55	18	Female	
6	69	1.50	89	Female	
7	71	1.78	42	Female	
8	64	1.67	16	Female	
9	69	1.64	52	Female	

Next steps: [Generate code with data](#) [New interactive sheet](#)

Convert Gender

```
data['Gender'] = data['Gender'].replace({
    'Female': 0,
    'Male': 1
})

data
```

/tmp/ipykernel_1273/3578139063.py:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version of pandas. To suppress this warning, please call `data['Gender'] = data['Gender'].replace({`

	Weight	Height	Age	Gender	
0	79	1.80	35	1	 
1	69	1.68	39	1	
2	73	1.82	25	1	
3	95	1.70	60	1	
4	82	1.87	27	1	
5	55	1.55	18	0	
6	69	1.50	89	0	
7	71	1.78	42	0	
8	64	1.67	16	0	
9	69	1.64	52	0	

Next steps: [Generate code with data](#) [New interactive sheet](#)

Check shape

```
data.shape

(10, 4)
```

Check null values

```
data.isnull().sum()

0
Weight 0
Height 0
Age 0
Gender 0

dtype: int64
```

Separate X and Y

```
x = data.iloc[:, 1:].values
y = data.iloc[:, :1].values

print(x)
print(y)
```

```
[[ 1.8  35.  1. ]
 [ 1.68 39.  1. ]
 [ 1.82 25.  1. ]
 [ 1.7  60.  1. ]
 [ 1.87 27.  1. ]
 [ 1.55 18.  0. ]
 [ 1.5  89.  0. ]
 [ 1.78 42.  0. ]
 [ 1.67 16.  0. ]
 [ 1.64 52.  0. ]]
[[79]
 [69]
 [73]
 [95]
 [82]
 [55]
 [69]
 [71]
 [64]
 [69]]
```

Split training and testing dat

```
x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2, random_state=42
)
```

create Linear Regression model

```
model = LinearRegression()

model.fit(x_train, y_train)

y_pred = model.predict(x_test)

y_pred
```

```
array([[56.85651025],
       [79.10576115]])
```

Print actual values

```
print(y_test)
```

```
[[64]
 [69]]
```

Print predicted values

```
print(y_pred)
```

```
[[56.85651025]
 [79.10576115]]
```

Check model score

```
model.score(x_test, y_test)
```

```
-11.252468343750825
```

Training graph

Because this is multiple linear regression, we can show Height against Weight while keeping Age and Gender fixed

```
import numpy as np

x_train_height = x_train[:, 0]
```

```

plt.scatter(
    x_train_height,
    y_train,
    color='blue',
    label='Actual Weight'
)

mean_age = x_train[:, 1].mean()

mode_gender = pd.Series(
    x_train[:, 2].astype(int)
).mode()[0]

height_range = np.linspace(
    x_train_height.min(),
    x_train_height.max(),
    100
)

regression_data = np.array([
    [h, mean_age, mode_gender]
    for h in height_range
])


predicted_weight = model.predict(regression_data)

plt.plot(
    height_range,
    predicted_weight,
    color='red',
    label='Regression Line'
)

plt.title('Weight vs Height (Training Set)')
plt.xlabel('Height (meters)')
plt.ylabel('Weight (kg)')
plt.legend()
plt.grid(True)

plt.show()

```



Testing graph

```

x_test_height = x_test[:, 0]

plt.scatter(
    x_test_height,
    y_test,
    color='blue',
    label='Actual Weight'
)

height_range_test = np.linspace(
    x_test_height.min(),
    x_test_height.max(),

```

```
100
)

regression_data_test = np.array([
    [h, mean_age, mode_gender]
    for h in height_range_test
])

predicted_weight_test = model.predict(
    regression_data_test
)

plt.plot(
    height_range_test,
    predicted_weight_test,
    color='red',
    label='Regression Line'
)

plt.title('Weight vs Height (Testing Set)')
plt.xlabel('Height (meters)')
plt.ylabel('Weight (kg)')
plt.legend()
plt.grid(True)

plt.show()
```

